

Advisory Text Is Not Ground Truth: Reporter Coverage, Campaign Duplication, and a Null Result for Package Naming in GHSA/OSV Malware Advisories

Tahsin Tajwar Tanni and Nafiz Khan
Independent Researchers
Dhaka, Bangladesh

Abstract—Advisory databases such as GHSA and OSV are increasingly mined as ground truth for malware behaviour in package ecosystems. We collect 401 confirmed-malicious PyPI and npm advisory records and show that they correspond to only 200 packages and 94 independent campaigns once cross-source re-reporting, multi-name campaigns and reporter-assigned campaign identifiers are collapsed. We then test whether packages with compound/generic names (e.g. `py-*`, `*-tools`, `*-client`) are more or less likely to have a described payload mechanism. Pooled, the odds ratio (OR) is 0.88 [0.35, 2.21]; adjusted for ecosystem and reporter coverage, it is 0.91 [0.33, 2.52]. The null is stable across model specifications and across three successive versions of the campaign grouping. Reporter coverage predicts the label far better than naming does: reporter count and Amazon Inspector coverage each carry an OR of 3.9–8.6 when entered alone. Reading only GHSA’s text instead of the GHSA+OSV union moves the estimate to OR 2.85 ($p = 0.047$), because GHSA omits 21 mechanism descriptions that OSV carries. A human check of 51 labeler positives found precision 0.94 [0.83, 0.98]. The naming grammar, originally motivated by LLM package hallucination, flags only 5% of 139 real hallucinated names versus 16–21% of popular benign packages, so it does not measure hallucination-proneness. We release the frozen corpus, the campaign-collapse method and all analysis code.

Index Terms—software supply chain, malicious packages, security advisories, PyPI, npm, measurement validity

I. INTRODUCTION

Malicious packages account for 49% of vulnerability reports in npm and 14% in PyPI [10], and LLM code assistants add a new route to them: a model recommends a package that does not exist, and an attacker registers the name [1], [4]. Measurement work on this threat increasingly treats the GitHub Security Advisory database (GHSA) and OSV as ground truth, both for which packages are malicious and for what the malware does.

This paper checks two assumptions that such work usually makes and rarely tests. First, that advisory records are independent observations. They are not: GHSA and OSV re-report the same package under different identifiers, and one attacker squats many names with one payload. In our corpus, 401 records reduce to 94 campaigns. Second, that a mechanism mentioned in the advisory text is a property of the malware. It is partly a property of whoever wrote the advisory: which

analysts covered a package changes whether a mechanism is described at all. Once both are accounted for, the association we set out to test disappears. The null result is the consequence of measuring carefully, and we report it together with the two corrections because both change how the next study should be designed.

Research question. Among confirmed-malicious packages, is compound/generic naming associated with a described payload mechanism, after accounting for how the advisory was reported?

Contributions.

- 1) A frozen, checksummed, campaign-level corpus (401 records $\rightarrow$ 200 packages $\rightarrow$ 94 campaigns) and a reusable two-stage collapse.
- 2) Evidence that mechanism labels mined from advisory text depend strongly on reporter coverage.
- 3) A stable null for naming $\rightarrow$ mechanism, with a power analysis.
- 4) A negative validation of a compound-naming grammar as a proxy for LLM hallucination.

II. BACKGROUND AND RELATED WORK

Package hallucination. Spracklen et al. [1] measured hallucinated package names across 16 code-generating models and found the problem persistent across models. Churilov [8] replicated the study on 2026 frontier models and released the set of names hallucinated by all five models tested, which we use in Section III-C. Huang et al. [4] registered hallucinated npm names and observed real installs, which establishes the threat model behind this work.

Malicious-package measurement. Akhavani et al. [10] analyse vulnerability reports across ten language ecosystems and find that malicious packages make up a large share of reported npm and PyPI vulnerabilities. Guo et al. [9] benchmark npm malicious-package detectors on a large labelled set. Hillah et al. [3] use PyPI registration metadata to give a calibrated posterior for registered-but-suspicious packages. Guo et al. note that typosquatting campaigns reuse identical payloads across many names, but treat each package as a separate instance; Akhavani et al. remove 6,391 duplicate reports and

prefer Snyk text over GitHub advisories because it is more detailed, a reporter-coverage choice whose effect on estimates they do not measure. We quantify both.

Scope. Work on hallucinated agent-skill names that host adversarial prompts [2] concerns a different artifact and a different mechanism, and is not studied here.

III. DATA AND METHOD

Fig. 1 shows the pipeline.

A. Collection

We query the GHSA API for advisories of type `malware` in the pip and npm ecosystems and query OSV.dev for each package found. After removing exact duplicates on (ecosystem, package, advisory id), the corpus has 401 records for 200 packages (100 per ecosystem). GHSA is a live API and returns advisories newest first. We fetched five pages, so 386 of the 401 records were published between 1 September and 2 October 2026. The pull is frozen to a snapshot with a SHA-256 manifest, and every number in this paper is computed from that file.

B. Campaign collapse

Fisher’s and CMH tests assume independent units; without collapsing, one npm table reached an OR of 178 on single-digit cells.

We collapse in two stages. Stage 1 merges records of the same (ecosystem, package), which catches GHSA/OSV re-reports. Stage 2 groups packages by a campaign signature, checked in priority order:

- 1) the reporter’s own `Campaign: identifier`, when present (the kam193 reporter assigns these);
- 2) otherwise the first 200 characters of the description, after stripping `## Source: headers` and masking the package’s own name.

Generic reporter templates (GHSA’s “should be considered fully compromised” text and kam193’s generic install-time pentest template) are never used as keys, because they are identical across unrelated incidents; packages carrying only such text stay as their own unit. A unit counts as grammar-flagged, or as DV-positive, if any member does. Choosing one representative record instead silently moves campaigns between table cells, because names within one campaign can differ in their grammar flag.

The result is 200 packages $\rightarrow$ 94 campaigns (npm 22, PyPI 72). The npm side is dominated by a few families: two WhatsApp-library fork families (45 and 15 names), a 10-name dependency-confusion family, and two related families of 10 and 3 names. Together these hold 83 of the 100 npm packages.

The human labeler check (Section III-D) found problems before it produced a precision figure. Reading the sampled rows surfaced two grouping failures: per-package `## Source: headers` made identical texts look different, and a 10-name family whose texts begin with the package name was never merged. These led to the header strip and the name mask. It also surfaced a negation false positive, “it does not exfiltrate credentials ...”, which led to the negation guard.

TABLE I
GRAMMAR FLAG RATE BY SET, WITH WILSON 95% CIs.

Set	n	Flagged	Rate [95% CI]
Hallucinated names [8]	139	7	0.050 [0.025, 0.100]
Benign top-5000 PyPI	5,000	821	0.164 [0.154, 0.175]
Benign top-5000 npm	5,000	1,031	0.206 [0.195, 0.218]
Malware campaigns	94	25	0.266 [0.187, 0.363]

C. Independent variable: compound/generic naming

A name is flagged if it contains a token from lists of compound suffixes, compound prefixes or trend suffixes (e.g. `py-`, `-tools`, `-client`, `-sdk`). The classifier reads the name only. The grammar was designed as a proxy for names an LLM might hallucinate, so we tested it against real hallucinations (Table I). It flags 7 of the 139 names that all five models in Churilov’s study [8] hallucinated, against 16–21% of the top-5,000 benign packages. We therefore describe the variable as compound/generic naming, which is what it measures. The benign sets must be chosen with care: an npm control drawn from search-API results for common terms is flagged at 73%, because the search terms are themselves grammar tokens. The seven flagged hallucinated names are `aws-cdk`, `simics-6-api`, `flask-idempotency`, `aws-retry`, `rollup-plugin-es6`, `react-randomized` and `iana-language-tag`.

D. Dependent variable: described mechanism

A first taxonomy, aimed at agent-directed prompt injection, found nothing: advisory text is third-person analyst narration, not text from the package. The dependent variable is therefore whether the advisory text describes at least one of six mechanisms:

- obfuscated payload execution;
- a hidden install hook;
- credential or wallet exfiltration;
- remote backdoor access;
- remote payload retrieval;
- a brand-impersonation narrative.

The patterns were first derived from PyPI text and flagged 0 of about 200 npm records across four pipeline runs. Adding npm phrasing (`postinstall` hooks, JavaScript obfuscators, verbs such as *reads* and *POSTs*) fixed this. A negation guard drops a match when a cue (*does not*, *do not*, *did not*, *doesn’t*, *no evidence of*) occurs within 40 characters before it, unless *but*, a semicolon or a period intervenes. The labeler’s signature accepts text only, never the name, and a test enforces this.

The text is the union of a package’s GHSA and OSV descriptions. Under this setting, 21 packages are positive only because of text that OSV carries and GHSA omits: 20 through an amazon-inspector section and one through a kam193 section. On reading, all 21 describe real mechanisms. GHSA-only text is reported as a sensitivity analysis (Section IV-D).

One author annotated 51 labeler positives under a written rule: all 21 OSV-only positives, plus 30 sampled campaigns




Fig. 1. Data collection, labeling, and analysis pipeline. The DV and IV labelers in step 3 read disjoint inputs, so the association test is not circular.

(an npm census of 12 and a seeded draw of 18 PyPI). Matches found only in analyst tag lists, not in prose, count as FALSE. The result was 44 REAL, 3 FALSE and 4 UNSURE, giving precision 0.936 [0.828, 0.978] excluding UNSURE and 0.863 [0.743, 0.932] counting UNSURE as FALSE. All 3 FALSE rows were tag-line-only matches, and 2 of them contain prose describing the behaviour that the regex missed, so as a measure of label correctness these figures are conservative. Annotators saw package names, and recall was not measured.

E. Reporter coverage

We parse the `## Source: headers` in both texts of every package and merge them per campaign. GHSA’s boilerplate and OSV’s `ghsa-malware` section are the same text and count as one reporter. Headerless text without boilerplate is *unattributed* (73 of 200 packages). From the parsed headers we derive a reporter count (`n_reporters`) and three indicators: `kam193`, `amazon-inspector` and `boilerplate`.

F. Analysis

We use Fisher’s exact test on the pooled and per-ecosystem 2×2 tables, and CMH tests stratified by ecosystem and by ecosystem $\times$ reporter count [6]. We fit the logistic regression $DV \sim grammar + ecosystem + n_reporters + amazon_inspector$ (model d), and three variants: (a) without the Amazon term, (b) without the reporter count, and (c) (a) plus the boilerplate and `kam193` indicators. ORs come with Woolf, Robins–Breslow–Greenland or Wald 95% CIs as appropriate. Power at the achieved n follows Hsieh et al. [7].

IV. RESULTS

A. Descriptive

Of the 94 campaigns, 25 are grammar-flagged and 51 are DV-positive. The DV rate depends strongly on who reported the campaign:

- `amazon-inspector` coverage: 0.65 (46/71), against 0.22 (5/23) without;
- `kam193` coverage: 0.59 (37/63);
- boilerplate among the reporters: 0.38 (5/13); boilerplate as sole reporter: 0/4.

Reporter coverage is not independent of naming either. The GHSA boilerplate appears in 0 of 74 records of flagged packages, against 26 of 327 records of unflagged packages.

B. Main test

The pooled table is $[[13, 12], [38, 31]]$, which gives OR 0.884 [0.353, 2.211], $p = 0.818$ (Table II). Adjusting for ecosystem with CMH gives 0.890 [0.355, 2.232], $p = 0.802$. The npm stratum ($n = 22$) is too small to interpret on its own.

C. Adjusted

In model (d) the grammar OR is 0.912 [0.331, 2.515], $p = 0.858$. Reporter coverage explains the label instead. Entered without its collinear partner, the reporter count has OR 3.90 (model a) and Amazon Inspector coverage OR 5.86 (model b), both $p \leq 0.003$. With both terms in the model (variance inflation factors 2.3), neither is individually significant, but model (d) as a whole is (likelihood-ratio $p = 0.0036$). The grammar flag’s variance inflation factor is 1.01.

D. Text-source sensitivity

Reading only GHSA’s text gives OR 2.851 [1.107, 7.341], $p = 0.047$ (CMH 2.907, $p = 0.023$). We do not treat this as a finding, for four reasons:

- It is one of many text, grouping and signature settings we examined, with no correction for multiple comparisons.
- Its npm stratum has a zero cell: none of the 16 unflagged npm campaigns is DV-positive under GHSA text.
- The regression with boilerplate and `kam193` terms does not converge, because no boilerplate campaign is DV-positive under GHSA text.
- The gap is explained by reporter coverage. None of the 21 packages that are positive only through OSV-carried text is grammar-flagged, against 37 of the other 179 packages (Fisher $p = 0.016$).

A single-source estimate is therefore not “cleaner”; it is biased by which reporter’s text was dropped.

E. Power

With 26.6% of campaigns flagged and a baseline DV rate of 0.55 among unflagged campaigns, the design has 0.31 power at OR 2.0 and 0.65 at OR 3.0. Reaching 80% power at OR 2.0 needs about 339 campaigns.

V. DISCUSSION

What the null means. We find no evidence that compound/generic naming goes with a described mechanism. At $n = 94$ this rules out only large effects (OR $\gtrsim 3$); it is

TABLE II
GRAMMAR FLAG VS. DESCRIBED MECHANISM, 94 CAMPAIGNS, GHSA+OSV TEXT UNLESS NOTED. REGRESSION MODELS (A)–(D) AS DEFINED IN SECTION III-F.

Specification	n	Grammar OR [95% CI]	p	Reporter terms, OR [95% CI]	p
Pooled, Fisher	94	0.884 [0.353, 2.211]	0.818		
npm, Fisher	22	2.200 [0.323, 14.975]	0.624		
PyPI, Fisher	72	0.673 [0.234, 1.940]	0.587		
CMH by ecosystem	94	0.890 [0.355, 2.232]	0.802		
CMH by ecosystem $\times$ reporter count	94	1.067 [0.402, 2.830]	0.891		
Logit (d)	94	0.912 [0.331, 2.515]	0.858	n_reporters 1.862 [0.528, 6.568]	0.334
				amazon-inspector 3.388 [0.701, 16.372]	0.129
Logit (a)	94	0.992 [0.360, 2.732]	0.987	n_reporters 3.898 [1.572, 9.664]	0.003
Logit (b)	94	0.855 [0.317, 2.310]	0.758	amazon-inspector 5.857 [1.898, 18.073]	0.002
Logit (c)	94	0.827 [0.279, 2.448]	0.731	n_reporters 8.623 [2.316, 32.098]	0.001
				boilerplate 0.211 [0.022, 2.070]	0.182
				kam193 0.158 [0.024, 1.030]	0.054
<i>Sensitivity: GHSA text only, Fisher</i>	94	2.851 [1.107, 7.341]	0.047		

not evidence that no effect exists. A cost-minimising attacker model [5] could still motivate a directional hypothesis for a larger study, but it cannot explain a result we did not observe.

What the reporter finding means for others. Any study that uses GHSA or OSV text as a behavioural label inherits a reporter-coverage confound. Dropping one source changed our estimate from 0.88 to 2.85, because the dropped reporter’s detail went disproportionately to unflagged packages. We recommend three practices:

- use the union of the source texts;
- parse reporter sections;
- control for reporter count.

Why we do not claim “sophistication”. The dependent variable is whether any mechanism is described. A count of distinct categories shows no difference either: a mean of 0.76 for flagged and 0.70 for unflagged campaigns (Mann–Whitney $p = 0.92$).

Why we do not claim “hallucination”. The grammar flags 5% of real hallucinations, most of which are import or module names used as package names (`objc`, `opentelemetry`, `tencentcloud`), so it measures generic naming. The gradient from 16–21% (benign) to 27% (malware campaigns) is descriptive, and the sets are not matched.

Ecosystems. The npm stratum consists of a handful of families plus singletons, so this study supports no per-ecosystem claim.

VI. THREATS TO VALIDITY

Construct. The labeler is a regex taxonomy. Its check is a single-annotator review of 51 positives (precision 0.94 [0.83, 0.98]), with no inter-annotator agreement and no recall estimate. The grammar is validated only as a measure of generic naming.

Internal. Reporter coverage is modelled but not exhausted: 37% of packages are unattributed. As superseded history, the campaign grouping changed three times during the study (142 $\rightarrow$ 106 $\rightarrow$ 94 campaigns) as the labeler check surfaced merge failures, and further under-merging is possible. The point

estimate crossed 1 across those changes (0.85, 1.18, 0.88), which itself suggests there is no signal at this n .

External. The corpus covers a five-week window, two ecosystems and 94 campaigns. Of these, 22 are npm campaigns dominated by a few families.

Statistical. The study is underpowered for $OR < 3$, several strata are sparse, and CMH assumes a homogeneous OR across strata.

VII. CONCLUSION AND FUTURE WORK

Advisory records are not independent, and advisory text is not ground truth. After correcting for both, compound/generic naming shows no association with described malware mechanisms, while reporter coverage does. Next steps:

- grow the corpus backward in time to about 350 campaigns;
- validate the labeler with two annotators and add a matched benign comparison;
- build a hallucination instrument based on import names rather than naming grammar.

Artifact availability. The frozen snapshot, manifests, analysis code and tests are available at [anonymized]. GHSA advisory data are licensed CC BY 4.0, OSV records carry per-source licences, and the hallucinated-name set [8] is CC BY 4.0.

ACKNOWLEDGMENTS

The authors used an AI assistant (Claude, Anthropic) for help with code and drafting. The authors reviewed all code, analyses and text and take full responsibility for the content of this paper.

REFERENCES

- [1] J. Spracklen, R. Wijewickrama, A. H. M. N. Sakib, A. Maiti, B. Viswanath, and M. Jadhwal, “We have a package for you! A comprehensive analysis of package hallucinations by code generating LLMs,” in *Proc. 34th USENIX Security Symposium*, 2025, arXiv:2406.10279.

- [2] A. Spira, S. Cohen, E. Feldman, R. Bitton, A. Wool, and B. Nassi, "Beware of agentic botnets: Scalable untargeted promptware attacks via universal and transferable adversarial HalluSquatting," *arXiv preprint arXiv:2607.07433*, 2026.
- [3] L. M. Hillah, J.-M. Richard, and R. Hasnaoui, "Bayesian-calibrated detection of hallucinated package imports in AI-assisted code," *arXiv preprint arXiv:2606.13918*, 2026.
- [4] Z. Huang, C. Liu, Z. Wang, X. Cui, X. Deng, and M. Xu, "Package hallucinations as phantoms in open-source software supply chains: An empirical security analysis," *ACM Trans. Softw. Eng. Methodol.*, 2026.
- [5] L. Allodi, F. Massacci, and J. Williams, "The work-averse cyberattacker model: Theory and evidence from two million attack signatures," *Risk Analysis*, vol. 42, no. 8, pp. 1623–1642, 2022.
- [6] A. Agresti, *Categorical Data Analysis*, 3rd ed. Hoboken, NJ, USA: Wiley, 2013.
- [7] F. Y. Hsieh, D. A. Bloch, and M. D. Larsen, "A simple method of sample size calculation for linear and logistic regression," *Statistics in Medicine*, vol. 17, no. 14, pp. 1623–1634, 1998.
- [8] A. Churilov, "The range shrinks, the threat remains: Re-evaluating LLM package hallucinations on the 2026 frontier-model cohort," *arXiv preprint arXiv:2605.17062*, 2026.
- [9] W. Guo, Z. Chen, Z. Xu, C. Liu, M. Kang, S. Song, C. Liu, Y. Xu, W. Sun, and Y. Liu, "How effective are NPM malicious package detectors? A large-scale empirical study," *arXiv preprint arXiv:2603.27549*, 2026.
- [10] S. A. Akhavan, B. Ousat, and A. Kharraz, "Open source, open threats? Investigating security challenges in open-source software," *arXiv preprint arXiv:2506.12995*, 2025.